

# COMP40370 Practical 3

## REGRESSION AND CLASSIFICATION

Prof. Tahar KECHADI

### Assignment Files

- `./Practical-03-A.pdf` Assignment questions (this file).
- `./auto-mpg.csv` Data file.

### Expected output files

- `./Practical-03.ipynb` Python notebook solutions.
- `./Practical-03.html` Python notebook in HTML format.

Requirements: Python 3.9+, pandas 1.3+, sklearn 0.24+.


## Part A: Regression Analysis

This practical work aims to use regression to predict a continuous-valued variable. The dataset for this practical describes the fuel consumption of some cars, the same dataset used in Practical 01. The dataset file is called “`auto-mpg.data`”.

### Question 1: Polynomial Regression

- according to the data pre-processing instructions in Practical 01:
  - Fill in the missing values, replace the outliers and remove duplicates rows of the dataset.
  - Replace the number of column ‘origin’ by the country of origin as follows: 1:‘USA’,2:‘Europe’,3:‘Japan’
  - Create a new column ‘brand’, and extract the car brand from ‘car name’ column, replace any misspelling as follows: ‘chevrolet’: ‘chevrolet’, ‘chevy’: ‘chevrolet’, ‘volkswagen’: ‘volkswagen’, ‘vw’: ‘volkswagen’, ‘hi’: ‘harvester’, ‘maxda’: ‘mazda’, ‘toyouta’: ‘toyota’, ‘mercedes-benz’: ‘mercedes’. Feel free to pre-process it as you see fit.
- Generate a simple linear regression model that predicts `mpg` based on ‘horsepower’ alone (use `sklearn.linear_model` library) and train it with 70% of the data. What are the RMSE and  $R^2$  values of testing/predicted data?
- Generate a group of polynomial models (use `sklearn.preprocessing.PolynomialFeatures`) that predict `mpg` based on the ‘horsepower’ alone. The models should have degrees of 2, 5, 10, and 20. For each model, find RMSE and  $R^2$  values of testing/predicted data.
- Plot a scatter diagram between `mpg` and ‘horsepower’ with all the fitted linear and non-linear models. Explain what will happen when the model complexity is increased. Which model is better? Explain your answer.

An **illustration** of scatter diagram:



## Question 2: Multiple Linear Regression

- Predict `mpg` based on 'horsepower', 'displacement', 'weight' and 'acceleration'. What are the RMSE and  $R^2$  values? Normalize the data then make the prediction again, What are the RMSE and  $R^2$  values after normalization?
- Which are the most influential two factors in the `mpg` prediction?
- Predict `mpg` based on the first PCA component of the above four variables ('horsepower', 'displacement', 'weight' and 'acceleration'). Do you see any accuracy reduction?
- Add origin and cylinders as categorical variables appropriately into the prediction model and discuss any accuracy changes.
- By adding "model\_year" as an ordinal variable and "brand" as a categorical variable do you see any improvement in the model performance?

## Part B: Classification – Decision Trees

- Create a new column from pre-processed data-frame in Q1 (a) called 'FEGroup' to categorise cars as "high-fuel" and "low-fuel" consumption. Allocate 10% of cars having the lowest mpg (use pandas `qcut` function) into the 'high-fuel' consumption class. What is the base accuracy? The same question for the case where all cars are classified as 'low-fuel'.
- Using `sklearn.tree` library, generate a decision tree based on 'horsepower', 'weight', 'acceleration', 'cylinders' and 'origin'. Using information gain (entropy) as the splitting criterion, set `max_depth` to 3 while leaving everything else to their default values.
- Find the training and test accuracies when the data is split into 70% for training and 30% for testing. (**Hint:** use `np.random.seed(42)`). Plot the resulting decision tree using `sklearn plot_tree` function.
- Train other decision trees, but this time with `max_depth` set to 3, 5, 8, respectively. Compare the results of these 3 DTs with the previous model. Discuss the problem of measuring accuracy in an imbalance class problem, such as this one.

**Please make sure that you have completed this practical. Next week, you will get the second part of the practical.**